

KOALA: Estimating coalition probabilities in multi-party electoral systems

Do you have a subtitle?
If so, write it here

Alexander Bauer · Andreas Bender ·
André Klima · Helmut Küchenhoff

Received: date / Accepted: date

Abstract 150 to 250 words

Common election poll reporting is often misleading as sample uncertainty is either not covered at all or only insufficiently. For a more comprehensive coverage, we propose shifting the focus towards reporting survey-based probabilities of specific election outcomes. We present such an approach for multi-party electoral systems, focusing on probabilities of coalition majorities. A Monte Carlo based Bayesian Multinomial-Dirichlet model is used for estimation. The method utilizes published opinion polls conducted by established polling agencies and is accompanied by a pooling approach to summarize multiple current surveys to reduce sample uncertainty, thereby accounting for dependencies between polling agencies. Potential biases of specific agencies are not taken into account. Sample uncertainty-based probabilities are estimated, assuming the election was held today, not accounting for potential shifts until election day. Possible visualizations of the results are shown and the benefit of the approach is outlined by discussing the German federal elections 2013 and 2017. An implementation is freely available in the R package **coalitions**.

Keywords 4 to 6 keywords Election analysis · Opinion polls · Election reporting · Multinomial-Dirichlet · Pooling

A. Bauer
Statistical Consulting Unit StaBLab, Department of Statistics, LMU Munich, Germany
Tel.: +49-89-2180-3197
Fax: +49-89-2180-5308
E-mail: alexander.bauer@stat.uni-muenchen.de
ORCID: 0000-0003-3495-5131

A. Bender
Statistical Consulting Unit StaBLab, Department of Statistics, LMU Munich, Germany
ORCID: 0000-0001-5628-8611

A. Klima
Statistical Consulting Unit StaBLab, Department of Statistics, LMU Munich, Germany

H. Küchenhoff
Statistical Consulting Unit StaBLab, Department of Statistics, LMU Munich, Germany

1 Introduction and data

Election polls are conducted by different polling agencies and try to represent the public opinion based on a finite sample. Usually, polling agencies publish the shares of the electorate who would vote for the reported political parties *if the election was held today*, the number of overall respondents and – more or less prominent – information about the uncertainty of the results. Current reporting of general media on such surveys however is most often limited to the observed shares, while sample uncertainty is usually ignored or only covered insufficiently.

Many examples for inaccurate reporting can be found in multi-party electoral systems. In such, one party usually doesn't obtain enough votes for an outright majority. In these situations, multiple parties form a so-called coalition to jointly obtain the necessary majority of seats in parliament. Oftentimes, a coalition in the media is stated to "lose" its majority just because the joint poll share drops under 50% from one opinion poll to the next (cf. FAZ.net, 2017). Such interpretations are clearly misleading as election polls are based on a finite sample of voters and only allow conclusions about the whole electorate with a specific certainty. Reporting results in this manner reinforces general misunderstandings of the public about which message to draw from such surveys.

Beyond ensuring proper reporting of sample uncertainties, in our opinion, the whole focus in election poll reporting should be shifted from describing the raw observed party shares. Instead, reporting should focus on the most relevant question, i.e. *how probable* specific events or election outcomes are. As probabilities combine both the observed raw voter share and sample uncertainty in one number, they are generally easier to communicate to the general public and less at risk to be misinterpreted. Events of interest can range from probabilities of parties passing a specific voter share via one party obtaining more seats in parliament than another through to probabilities of majorities of potential multi-party coalitions. We present our KOALA (Coalition Analysis) approach to estimate such probabilities to bring more value to opinion poll-based reporting, specifically focusing on multi-party electoral systems and the estimation of probabilities of coalition majorities. To estimate the probabilities, a Bayesian Multinomial-Dirichlet model with Monte Carlo simulations is used. Also, a pooling approach is presented to summarize multiple current opinion polls to reduce sample uncertainty. Prior to the German federal elections 2013 and 2017, results based on (an earlier iteration of) our approach already entered general media reporting (cf. ZEIT ONLINE, 2013; Gelitz, 2017).

As database, we use opinion polls conducted by established polling agencies, quantifying the electoral behavior *if an election was held today*. We focus on the question of quantifying current majority situations, not taking into consideration potential shifts until election day. Approaches for predicting future election outcomes based on past information can e.g. be found in Graefe (2017) or Norpoth and Gschwend (2010).

All methods are implemented in R (R Core Team, 2017) and are available in the open-source package `coalitions` on GitHub (Bender and Bauer, 2018). An interactive shiny-based (Chang et al., 2017) website `koala.stat.uni-muenchen.de` visualizes estimated coalition probabilities and is used to communicate the results for German federal and state-wide elections to the general public. The process of fetching new polls, updating the website and sending out Twitter messages based

on the newest results is automated and allows for an immediate transfer of the estimated event probabilities to media and pulic.

Absatz zu: In Section 2 we will discuss ..., in 3 ...

2 Calculation of probabilities

In the last opinion poll conducted before the German federal election 2013 (Forsa, 2013), special interest was on whether CDU/CSU-FDP (also “Union-FDP”) would obtain enough votes to form the governing coalition:

Table 1 Observed voter shares in the Forsa opinion poll for the German federal election, published September 20th, 2013 with $n = 1995$ respondents

Union	SPD	Greens	FDP	The Left	Pirates	AfD	Others
40%	26%	10%	5%	9%	2%	4%	4%

The German election system mandates a 5% votes share for parties to enter the parliament. Votes for parties below this threshold are redistributed (proportionally) to parties above it. Redistribution of the upper poll leads to the following party shares:

Table 2 Redistributed party shares based on the Forsa opinion poll for the German federal election, published September 20th, 2013 with $n = 1995$ respondents. Parties marked with “–” didn’t pass the 5% hurdle.

Union	SPD	Greens	FDP	The Left	Pirates	AfD	Others
44.44%	28.89%	11.11%	5.56%	10.00%	–	–	–

As can be seen in Table 2, Union-FDP with its 45% raw voter share would get exactly 50% of parliament seats after redistribution. Thus, ingoring uncertainty one would conclude that a majority of the coalition is slightly missed. However, it is clear that this only holds with a certain probability and particularly depends on whether FDP and/or AfD pass the 5% hurdle.

The model to estimate event probabilities is based on the Multinomial distribution of the observed number of voters per party. We use a Bayesian Multinomial-Dirichlet model with Jeffreys prior as an uninformative prior for the true party shares θ_j (Gelman et al., 2013):

$$\begin{aligned} \theta &= (\theta_1, \dots, \theta_k)^T \sim \text{Dirichlet}(\alpha_1, \dots, \alpha_k), \\ \text{with } \alpha_1 &= \dots = \alpha_k = \frac{1}{2} \end{aligned} \tag{1}$$

Given one (pooled) survey, the posterior in this case also is a Dirichlet distribution with $\alpha_j = x_j + \frac{1}{2}$ for each party j and its observed vote counts x_j .

Using Monte Carlo simulations of election outcomes, one can obtain specific event probabilities by taking their relative frequency of occurrence. E.g., Union-FDP got a seat majority in 2 633 of 10 000 simulations, which equals an estimated probability of 26.33%.

If it is known that specific parties are underrepresented or overrepresented in election polls, this information can be included in the model by using an informative prior distribution. Underrepresented and overrepresented parties then would get higher and lower α_j values, respectively. However, such biases of polling agencies are hard to quantify as the true party share in the electorate is only known on election days. For our analyses, we therefore use an uninformative prior.

As voter shares are usually rounded before opinion polls get published, we adjust the available raw data by adding uniformly distributed random noise to the observed voter shares x_j in order to avoid potential biases caused by the use of rounded numbers:

$$\begin{aligned} x_{j,adj} &= x_j + r_{\gamma,j}, \\ \text{with } r_{\gamma,j} &\sim U[-\gamma, \gamma]. \end{aligned} \quad (2)$$

The correction coefficient γ is chosen according to rounding accuracy. E.g., for data rounded to 1% steps we use $\gamma = 0.5\%$. After random noise was added the adjusted shares are rescaled to ensure a sum of 100%. Overall, instead of using rounded numbers and simulating n_s values from the resulting posterior, we perform n_s simulations where we first adjust the voter shares using individually drawn random noise and then simulate one observation from the posterior.

Figure 1 shows the simulated parliament seat shares for the coalition Union-FDP, based on the observed voter shares in Table ?? . The estimated density is clearly bimodal as the observed FDP share before redistribution equals exactly 5% and so FDP only enters the parliament in half of the simulations. As mentioned, the estimated probability with 26.33% lies around a quarter.




Fig. 1 Density of 10 000 simulated parliament seat shares for the coalition Union-FDP before the German federal election in September 2013 based on the Forsa opinion poll in Table 1. The part of the density encoding for seat majorities is colored blue and resembles a probability of 26.33%.

As such density plots depict both the probability and the underlying uncertainty for specific coalitions, they are a nice possibility to communicate uncertainty underlying opinion polls to the general public.

To visualize the *development* of such probabilities for a specific coalition we recommend extending the visualization of Figure 1 by using ridgeline plots (Wilke, 2017) for the simulated seat distributions. In Figure 2 this plot type and the development of majority probabilities is compared to the redistributed raw shares, which are usually reported in media.

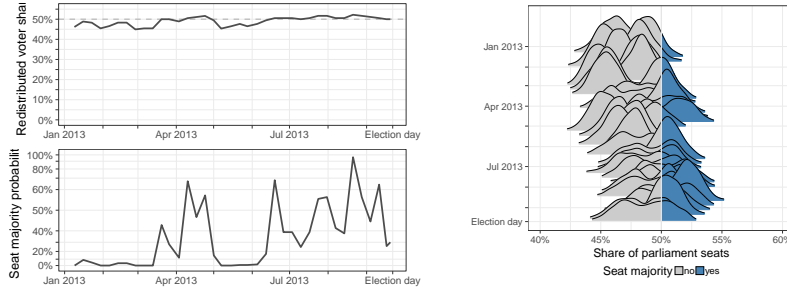


Fig. 2 Development of the prospect to form a government of the coalition Union-FDP before the German federal election in September 2013 based on Forsa opinion polls. Left top: Observed voter shares after redistribution. Left bottom: Probabilities to reach a majority of seats in parliament, based on 10 000 simulations. Right: Densities of simulated parliament seat shares based on 10 000 simulations. The parts of the densities encoding for seat majorities are colored blue.

Comparing the redistributed voter shares and the seat majority probabilities in Figure 2 it is clear that the probabilities consist of a lot more information than the voter shares. Especially, even small changes in the redistributed voter shares can make an immense difference regarding the chance of the coalition to form a government.

To focus on the most relevant changes in the majority probabilities, we propose the use of a skewed axis as shown in Figure 2. This axis stretches the range of values around 50% as a change in absolute probabilities from 40% to 60% is more relevant than a change from 70% to 90% as a success for achieving a seat majority is becoming more probable than a failure to do so, or in both cases a success is already very probable, respectively.

3 Pooling approach

In the presence of multiple published opinion polls, pooling is used to summarize the observed results in order to reduce sample uncertainty. To assure a reliable pooling regarding the current public opinion, we only use polls published within the past 14 days and only use the most recent survey published by each polling agency.

Looking at a single poll i , the observed number of votes X_{ij} for each of k parties follow a multinomial distribution with sample size n_i and underlying, unknown party shares θ_j in the population. Pooling over multiple such polls as independent random samples leads to another multinomial distribution for the summed number of votes $\sum_i X_{ij}$:

$$\sum_i X_{i1}, \dots, \sum_i X_{ik} \sim \text{Multinomial} \left(\sum_i n_i, \theta_1, \dots, \theta_k \right). \quad (3)$$

Further analyses, however, show that polls from different polling agencies are correlated and the independency assumption does not hold. Therefore, we adjust the resulting multinomial distribution by using an *effective sample size* (Hanley

et al., 2003), reflecting that the aggregation over multiple polls does not reflect the information from a sample with $\sum_i n_i$ observations.

Quantification of pairwise correlation is done based on the variance of the difference between two polls. The following equation holds for two independent random sample polls A and B :

$$\begin{aligned} \text{Var}(X_A - X_B) &= \text{Var}(X_A) + \text{Var}(X_B) - 2 \cdot \text{Cov}(X_A, X_B) \\ \Leftrightarrow \text{Cov}(X_{Aj}, X_{Bj}) &= \frac{1}{2} \cdot (\text{Var}(X_{Aj}) + \text{Var}(X_{Bj}) - \text{Var}(X_{Aj} - X_{Bj})). \end{aligned} \quad (4)$$

We take $\text{Var}(X_{Aj})$ and $\text{Var}(X_{Bj})$ as the theoretical variances of the binomially distributed, observed voter numbers and estimate $\text{Var}(X_{Aj} - X_{Bj})$ based on the observed differences between the party shares. Having done so, one can estimate the covariance $\text{Cov}(X_{Aj}, X_{Bj})$ and accordingly also the correlation. As the binomial distribution is directly proportional to the sample size, the effective sample size n_{eff} can be defined as the ratio between the estimated variance for the pooled sample and the theoretical variance of a sample of size one:

$$n_{\text{eff}} = \frac{\text{Var}(\text{pooled})}{\text{Var}(\text{sample of size 1})},$$

with, in the case of two surveys,

$$\text{Var}(\text{pooled}) = \text{Var}(X_A + X_B) = \text{Var}(X_A) + \text{Var}(X_B) + 2\text{Cov}(X_A, X_B)$$

and $\text{Var}(\text{sample of size 1})$ the theoretical variance of the pooled share.

Looking at the party-specific correlations between 20 surveys conducted by the two most regular German polling agencies, Emnid and Forsa, we on average end up with a medium high correlation, using mean party shares and sample sizes per institute for the theoretical variances. Other institute comparisons were not performed as too few published surveys were conducted cover comparable time frames. For simplicity, we do not recalculate the correlation for each simulation, but rather set the correlation used in our methodology to 0.5. For convenience, the calculation of n_{eff} is based on the party with most votes, as the specific party choice only marginally affects the results.

In the case of two published polls with 1500 and 2000 respondents, respectively, and a pooled share of 40% for the strongest party, the method leads to an effective sample size of $n_{\text{eff}} = 2341$. Thus, the method reduces sample uncertainty compared to using only single polls, but the method is also clearly conservative, compared to assuming independence between the polls and using a summed sample size of $1500 + 2000 = 3500$.

As noted, in practice we use a timeframe of 14 days, i.e. all surveys published at maximum 14 days ago are included in the calculation of the pooled sample. In the case of opinion polls for a specific election getting published only very rarely one can also extend this timeframe, e.g. using the heuristic of including all surveys published up to 14 days ago with weight 1 (using their whole sample size), and all surveys within 14 to 28 days before today with halved weight (using the halved sample size).

4 Application

We applied (an earlier iteration of) our method to the German federal elections 2013 and 2017, accompanied by news coverage in general media (cf. ZEIT ONLINE, 2013; Gelitz, 2017). In the following, we take these two elections as examples to underline the differences between standard media reporting on opinion polls and our approach, which focus on interpreting the raw observed shares after redistribution and the estimated Bayesian probabilities, respectively.

For each election, 2013 and 2017, we base the event probabilities on pooled surveys, which use the opinion polls published by the major German polling agencies, i.e. Allensbach, Emnid, Forsa, Forschungsgruppe Wahlen, GMS, Infratest dimap and INSA. Opinion poll data from these polling agencies is publicly available¹ on www.wahlrecht.de.

4.1 German federal election 2013

From 2009 to 2013, the German government was formed by a coalition of conservative CDU/CSU (or "Union") and the liberal FDP. Thus, main interest before the election on September 22nd was in whether this coalition could maintain its majority of parliament seats. As FDP in opinion polls was very close to a voter share of 5%, i.e. the minimum hurdle to pass into parliament, FDP successfully passing into parliament was crucial for the success of the coalition.




Fig. 3 Development of the pooled raw voter shares from October 2012 until election day on September 22nd, 2013, based on a pooling window of 14 days. As specific voter shares for AfD were only sparsely reported before the election in 2013 the party is contained in "Others".

Union-FDP coalition majority

Figure 4 visualizes the pooled opinion poll results for the potential coalition Union-FDP, since one year before the election. Starting from a redistributed voter share of around 43% in October 2012, the voter share in the pooled results rose steadily, reaching its maximum of ca. 51% about one month before election day.

Only taking this development of the redistributed voter share into consideration, one would conclude that – apart from a short time window in August/September of 2013 – a majority of the Union-FDP coalition would not be possible if elections would have been held. In the best case, sample uncertainty would be mentioned, making a majority for the coalition (*highly*) *improbable*, but a more solid quantification of uncertainty is not easily possible.

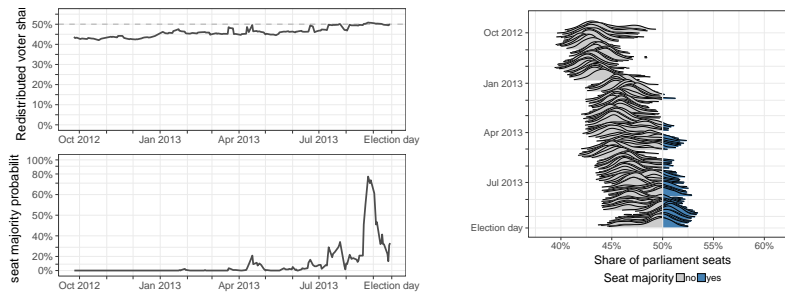


Fig. 4 Development of the prospect to form a government of the coalition Union-FDP before the German federal election in September 2013 based on pooled opinion polls. Left top: Observed voter shares after redistribution. Left bottom: Probabilities to reach a majority of seats in parliament, based on 10 000 simulations. Right: Densities of simulated parliament seat shares based on 10 000 simulations. The parts of the densities encoding for seat majorities are colored blue.

Using the calculated seat majority probabilities together with the simulated densities instead draws a more comprehensive picture of the situation. While the overall development of probabilities growing with time is the same as with the redistributed voter shares, there are two points that clearly get more weight when shifting the focus in reporting from a black-and-white reporting towards discussing *how probable* events are:

First of all, the impact of small changes in voter shares is nicely resembled in the probability development. Thus, a rise in voter share from 45% to 48% in March 2013 only has a very small impact on the absolute probability, whereas changes near a share of 50% are much more influential.

Second of all, probabilities do not only take the summed party shares into consideration, but also implicitly cover the uncertainty regarding e.g. whether FDP passes the 5% hurdle or not. Apart from the end of August 2013, FDP passing into the parliament is highly uncertain and thus majority probabilities of the coalition are rather small. This can best be seen in the ridgeline plot (right panel in Fig. 4), where the densities are bimodal if the raw voter share of FDP is close to 5% (see Fig. 5).

FDP passing the 5% hurdle

As mentioned, our approach can also be used to quantify the probability for other events than seat majorities, e.g. for the FDP successfully passing the 5% hurdle. Figure 5 again shows that small changes in raw voter share of the party dramatically influence the success probabilities. Overall, based on the opinion polls correctly representing the whole population, a success of the FDP seemed very uncertain in most times. Only around the end of August 2013 the pooled voter share rose to 6% and the corresponding probability rose to nearly 100%.




Fig. 5 Development of the prospect of FDP to pass the 5% hurdle before the German federal election in September 2013 based on pooled opinion polls. Left top: Observed voter shares before redistribution. Left bottom: Probabilities to pass the 5% hurdle, based on 10 000 simulations. Right: Densities of simulated raw voter shares based on 10 000 simulations. The parts of the densities encoding for successfully passing the hurdle are colored blue.

4.2 German federal election 2017

Similar to the previous election, prior to the election on September 24th of 2017 main interest lied in whether Union-FDP would succeed in getting a seat majority in parliament. However, the situation differed regarding that the FDP in opinion polls was safely over 5% (cf. Fig. 9). The same was true for the right-wing party AfD, which in 2013 slightly failed to getting into parliament. Also, another potential left-wing coalition between SPD, the green party and the Left was a valid option to a Union-reigned coalition.




Fig. 6 Development of the pooled raw voter shares from October 2016 until election day on September 24nd, 2017, based on a pooling window of 14 days.

Union-FDP coalition majority

The development of redistributed voter shares and majority probabilities for Union-FDP over the window of one year before election day is depicted in Figure 7. Especially as FDP was at all times over 5% in the pooled opinion polls, the development of voter shares is much more smooth, compared to the development of 2013. Again, it can be seen that even very small changes in voter share cause immense changes in probability. E.g., at the end of August 2017, the seat majority probability rises by 30% and reaches its maximum, only caused by a voter share rise of around 1%. Also, according to the ridgeline plot, redistributed voter shares of under 47% usually correspond to negligibly small majority probabilities.

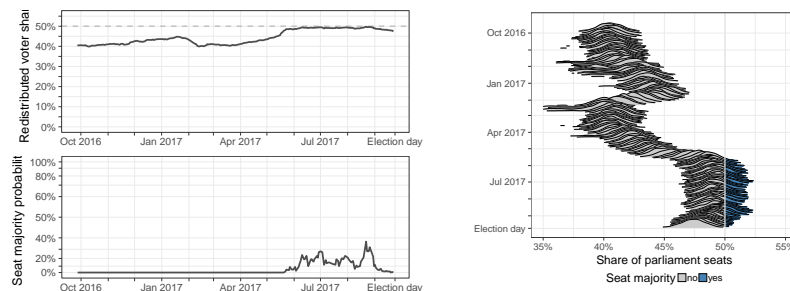


Fig. 7 Development of the prospect to form a government of the coalition Union-FDP before the German federal election in September 2017 based on pooled opinion polls. Left top: Observed voter shares after redistribution. Left bottom: Probabilities to reach a majority of seats in parliament, based on 10 000 simulations. Right: Densities of simulated parliament seat shares based on 10 000 simulations. The parts of the densities encoding for seat majorities are colored blue.

SPD-Left-Greens coalition majority

The same as for Union-FDP holds for the potential coalition of SPD, Greens and Left (see Fig. 8). Between February and April 2017 the redistributed voter shares reach a high plateau near 50%, and there a small change of about 1% at the end of March 2017 causes a rise of the probability of over 35%. In the end, especially SPD voter shares declined as heavily as they rose at the beginning of 2017, and thus a majority of the coalition shortly before election day was very unlikely.

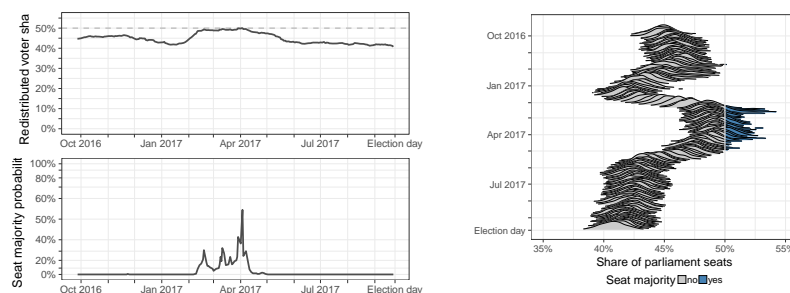


Fig. 8 Development of the prospect to form a government of the coalition SPD-Left-Greens before the German federal election in September 2017 based on pooled opinion polls. Left top: Observed voter shares after redistribution. Left bottom: Probabilities to reach a majority of seats in parliament, based on 10 000 simulations. Right: Densities of simulated parliament seat shares based on 10 000 simulations. The parts of the densities encoding for seat majorities are colored blue.

AfD becoming third strongest party

Specific interest in the election lied in which party would get the third best result after the major parties Union and SPD. In the year running up to the election, the right-wing party AfD oftentimes was leading in raw voter shares (see Fig. 9), but two weeks before election day three parties were almost head-to-head. Using our approach it is easily possible to estimate probabilities for the event that AfD becomes third biggest party in parliament, thereby quantifying the sample

uncertainty that underlies the opinion polls and summarizing this complex issue into one development of probabilities. The corresponding probability until April 2017 is always at least 70%, then varies heavily over the course of two months, drops to under 10%, and rises back up to over 90% shortly before election day, where the AfD is nearly 2% over its heaviest competitor – The Left – in absolute raw voter share.




Fig. 9 Development of the prospect of AfD to become the third strongest party before the German federal election in September 2017 based on pooled opinion polls. Left top: Observed voter shares before redistribution. Left bottom: Probabilities to become third strongest party, based on 10 000 simulations.

5 Conclusion

We presented a Bayesian approach to estimate probabilities of specific election outcomes based on publicly available opinion polls. Pooling allows for the inclusion of information from multiple surveys and reduces sample uncertainty. The estimated event probabilities are easy to communicate to the general public and prevent improper poll-based reporting as they readily include sample uncertainty. We visualize the results on a publicly available website for chosen elections and provide the open-source R package `coalitions` that allows for an easy application of the method to any multi-party electoral system. In this manner, our long-term goal is to make proper uncertainty assessment in general opinion poll-based reporting the rule, rather than an exception.

References

- Bender, A. and Bauer, A. (2018). *adibender/coalitions (Version v0.5.7)*. URL <http://doi.org/10.5281/zenodo.1172595>.
- Chang, W., Cheng, J., Allaire, J., Xie, Y., and McPherson, J. (2017). *shiny: Web Application Framework for R*. URL <https://CRAN.R-project.org/package=shiny>. R package version 1.0.5.
- FAZ.net (2017). Umfrage zur Bundestagswahl: Schwarz-Gelb verliert die Mehrheit. <http://archive.is/SuXVt>. Accessed 15 February 2018.
- Forsa (2013). Wenn am nächsten Sonntag Bundestagswahl wäre... <http://archive.is/f9vse>. Accessed 15 February 2018.

- Gelitz, C. (2017). Können die aktuellen Umfragen noch falschliegen? <http://archive.is/JydHd>. Accessed 15 February 2018.
- Gelman, A., Carlin, J. B., Stern, H. S., Dunson, D. B., Vehtari, A., and Rubin, D. B. (2013). *Bayesian data analysis*, volume 3. CRC press Boca Raton, FL.
- Graefe, A. (2017). The pollyvotes long-term forecast for the 2017 german federal election. *PS: Political Science & Politics*, **50**(3), 693–696.
- Hanley, J. A., Negassa, A., and Forrester, J. E. (2003). Statistical analysis of correlated data using generalized estimating equations: an orientation. *American journal of epidemiology*, **157**(4), 364–375.
- Norpoth, H. and Gschwend, T. (2010). The chancellor model: Forecasting german elections. *International Journal of Forecasting*, **26**(1), 42–53.
- R Core Team (2017). *R: A Language and Environment for Statistical Computing*. R Foundation for Statistical Computing, Vienna, Austria. URL <https://www.R-project.org/>.
- Wilke, C. O. (2017). *ggridges: Ridgeline Plots in 'ggplot2'*. URL <https://CRAN.R-project.org/package=ggridges>. R package version 0.4.1.
- ZEIT ONLINE (2013). Serie: Wahlistik. <http://archive.is/1SU1I>. Accessed 15 February 2018.